Enterprise AI Governance & Ethics Policy

Department: Compliance | Category: Governance | Enacted Year: 2024

Article 1: Core AI Principles and Human-in-the-Loop Oversight

All generative AI and machine learning systems deployed within enterprise operations must adhere to core ethical pillars: algorithmic fairness, transparency, auditability, data privacy, and human oversight. Autonomous decision-making is strictly prohibited in high-stakes scenarios such as employment screening, credit underwriting, and termination decisions.

A mandatory human-in-the-loop (HITL) review is required for any automated output that directly affects client financial status or employee welfare.

Article 2: Data Protection, PII Redaction, and Training Restrictions

Confidential corporate data, proprietary source code, and customer Personally Identifiable Information (PII) must be scrubbed using automated regex and NER token maskers prior to embedding generation or prompt transmission to third-party Large Language Model APIs.

Model inference logs and intermediate retrieval contexts must not be retained beyond a strict 90-day retention window. Enterprise customer data shall never be utilized to fine-tune external foundation models without explicit written bilateral consent.

Article 3: Evaluation Standards and Incident Reporting

AI pipelines undergo mandatory quarterly adversarial red-teaming. For Retrieval-Augmented Generation systems, hallucination rates must remain below a strict threshold of 2.0% on domain evaluation benchmark suites before production rollout.

Any discovered hallucination or security vulnerability categorized as Severity 1 (such as prompt injection exfiltrating PII) must be escalated to the AI Safety Committee and CISO within 24 hours of identification.